



Multiversion Concurrency Control



Multiversion Schemes

- Multiversion schemes keep old versions of data item to increase concurrency. Several variants:
 - **Multiversion Timestamp Ordering**
 - **Multiversion Two-Phase Locking**
 - **Snapshot isolation**
- Key ideas:
 - Each successful **write** results in the creation of a new version of the data item written.
 - Use timestamps to label versions.
 - When a **read**(Q) operation is issued, select an appropriate version of Q based on the timestamp of the transaction issuing the read request, and return the value of the selected version.
- **reads** never have to wait as an appropriate version is returned immediately.



Multiversion Timestamp Ordering

- Each data item Q has a sequence of versions $\langle Q_1, Q_2, \dots, Q_m \rangle$. Each version Q_k contains three data fields:
 - **Content** -- the value of version Q_k .
 - **W-timestamp**(Q_k) -- timestamp of the transaction that created (wrote) version Q_k .
 - **R-timestamp**(Q_k) -- largest timestamp of a transaction that successfully read version Q_k .



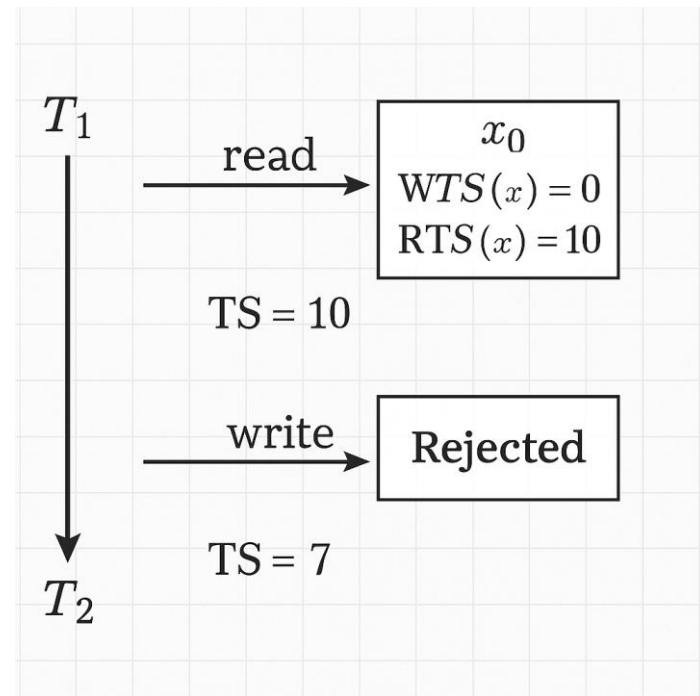
Multiversion Timestamp Ordering (Cont)

- Suppose that transaction T_i issues a **read**(Q) or **write**(Q) operation.
- Let Q_k denote the version of Q whose write timestamp is the largest write timestamp less than or equal to $TS(T_i)$.
 1. If **transaction** T_i issues a **read**(Q), then
 - the value returned is the content of version Q_k
 - If $R\text{-timestamp}(Q_k) < TS(T_i)$, set $R\text{-timestamp}(Q_k) = TS(T_i)$,
 2. If transaction T_i issues a **write**(Q)
 1. if $TS(T_i) < R\text{-timestamp}(Q_k)$, then transaction T_i is rolled back.
 2. if $TS(T_i) = W\text{-timestamp}(Q_k)$, the contents of Q_k are overwritten
 3. Otherwise, a new version Q_i of Q is created
 - $W\text{-timestamp}(Q_i)$ and $R\text{-timestamp}(Q_i)$ are initialized to $TS(T_i)$.



Multiversion Timestamp Ordering (Cont)

- Observations
 - Reads always succeed
 - A write by T_i is rejected if some other transaction T_j that (in the serialization order defined by the timestamp values) should read T_i 's write, has already read a version created by a transaction older than T_i .
- Protocol guarantees serializability



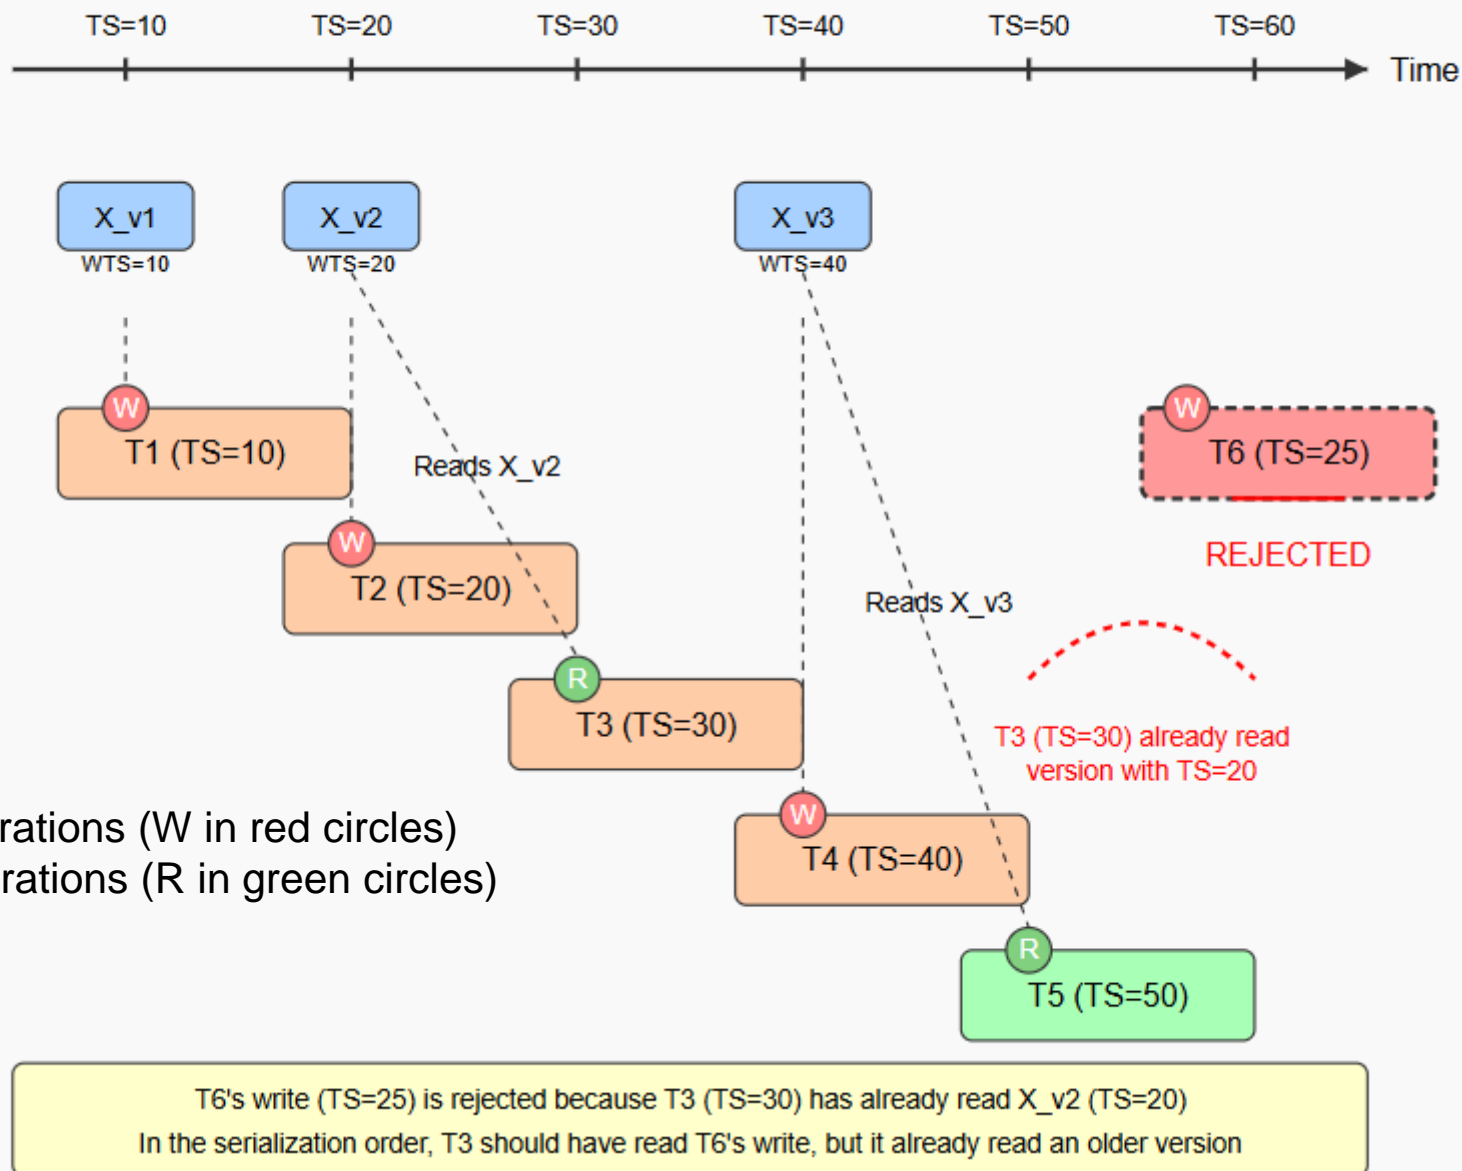


Multiversion Timestamp Ordering (Cont)

```
procedure MVTO_Write(x, T_i):  
    // Find the version x_j with the largest  
    WTS(x_j) < TS(T_i)  
    x_j ← latest_version_of(x) such that  
    WTS(x_j) ≤ TS(T_i)  
    if RTS(x_j) > TS(T_i) then  
        // Conflict: a future-timestamped  
        transaction already read the old value  
        reject T_i's write on x  
    else  
        // Safe: install new version  
        create version x_i with  
            WTS(x_i) ← TS(T_i)  
            RTS(x_i) ← TS(T_i)
```



Multiversion Timestamp Ordering



1. Write operations (W in red circles)
2. Read operations (R in green circles)



Multiversion Two-Phase Locking

- Differentiates between **read-only transactions** and update transactions
- **Update transactions** acquire read and write locks, and hold all locks up to the end of the transaction. That is, **update transactions follow rigorous two-phase locking.**
 - Read of a data item returns the latest version of the item
 - The first **write** of Q by T_i results in the creation of a new version Q_i of the data item Q written
 - $W\text{-timestamp}(Q_i)$ set to ∞ **initially to not allow other writes**
 - When **update** transaction T_i **completes, commit** processing occurs:
 - Value **ts-counter** stored in the database is used to assign timestamps
 - **ts-counter** is locked in two-phase manner
 - Set **$W\text{-timestamp}(Q_i) = (\text{ts-counter} + 1)$** for all versions Q_i that it creates
 - **$\text{ts-counter} = \text{ts-counter} + 1$**
 - Thereby, those transactions that start before T_i commits will see the value before the updates by T_i .



Multiversion Two-Phase Locking

Imagine an online store has a product **P** with a price history. Transactions update or read the price.

Initial State:

- The product **P** has an initial price: **\$100**, The **timestamp counter (ts-counter) = 10**.

Transactions:

Transaction T1 (Update Transaction)

- **Starts at ts = 11. Reads** the latest price (**\$100**).
- **First Write:** Creates a **new version P1** with price **\$120**.
 - The **write timestamp W-ts(P1)** is set to ∞ (blocking other writes).
- **Commit Process:**
- Locks **ts-counter**. Sets **W-ts(P1) = ts-counter + 1** \rightarrow **W-ts(P1) = 11**. Updates **ts-counter = 11**.

Transaction T2 (Read-Only Transaction)

- **Starts at ts = 10** (before T1 commits).
- **Reads product P**, but since **T1 hasn't committed**, it **sees the older price (\$100)**.



Multiversion Two-Phase Locking (Cont.)

■ Read-only transactions

- are assigned a **timestamp = ts-counter** when they start execution
- follow the multiversion timestamp-ordering protocol for performing reads
 - Do not obtain any locks
- Read-only transactions that start after T_i increments **ts-counter** will see the values updated by T_i .
- Read-only transactions that start before T_i increments the **ts-counter** will see the value before the updates by T_i .
- Only serializable schedules are produced.



Multiversion Two-Phase Locking (Cont.)

- Example
 - One product P with initial price: **\$100**
 - A **global timestamp counter (ts-counter) = 10**
- Transaction T1 (Update Transaction)**
- **Starts:** $ts(T1) = 11$
 - **Reads:** version P0 (price = \$100, W-ts = 10)
 - **Tentatively creates** a new version: P1 with price = \$120
 - $W\text{-}ts(P1) = \infty$ (not committed yet)
 - **At this moment:**
 - P0: price = \$100, $W\text{-}ts = 10$
 - P1: price = \$120, $W\text{-}ts = \infty$ (invisible to others)

Transaction	Start Time	Sees Version	W-ts of Version	Observed Price
T2 (read-only)	10	P0	10	\$100
T1 (update)	11	—	—	— (writes P1)
T3 (read-only)	11	P1	11	\$120



MVCC: Implementation Issues

- Creation of multiple versions increases **storage overhead**
 - **Extra tuples**
 - **Extra space** in each tuple for storing version information
- Versions can, however, be **garbage collected**
 - E.g., if Q has two versions Q5 and Q9, and the oldest active transaction has timestamp > 9 , then Q5 will never be required again
- Issues with
 - **primary** key and **foreign key constraint** checking
 - **Indexing** of records with multiple versions

See textbook for details



Snapshot Isolation

- Motivation: Decision support queries that read large amounts of data have concurrency conflicts with **OLTP** transactions that update a few rows
 - Poor performance results
- Solution 1: Use multiversion 2-phase locking
 - Give logical “snapshot” of database state to read only transaction
 - Reads performed on snapshot
 - **Update** (read-write) **transactions** use **normal locking**
 - **Works** well, but how does system know a transaction is **read only**?
- Solution 2 (partial): **Give snapshot of database state** to every transaction
 - Reads performed on snapshot
 - Use 2-phase locking on updated data items
 - Problem: variety of anomalies such as lost update can result
 - Better solution: snapshot isolation level (next slide)



Snapshot Isolation

- A transaction T1 executing with Snapshot Isolation
 - Takes **snapshot** of **committed** data at **start**
 - Always **reads/modifies data** in its **own snapshot**
 - **Updates** of **concurrent transactions** are not **visible** to **T1**
 - **Writes** of T1 **complete** when it **commits**
 - **First-committer-wins rule:**
 - ▶ Commits only if no other concurrent transaction has already written data that T1 intends to write.

T1	T2	T3
W(Y := 1) Commit		
	Start R(X) → 0 R(Y) → 1	
		W(X:=2) W(Z:=3) Commit
	R(Z) → 0 R(Y) → 1 W(X:=3) Commit-Req Abort	

Concurrent updates not visible
 Own updates are visible
 Not first-committer of X
 Serialization error, T2 is rolled back



Snapshot Read

- Concurrent updates invisible to snapshot read

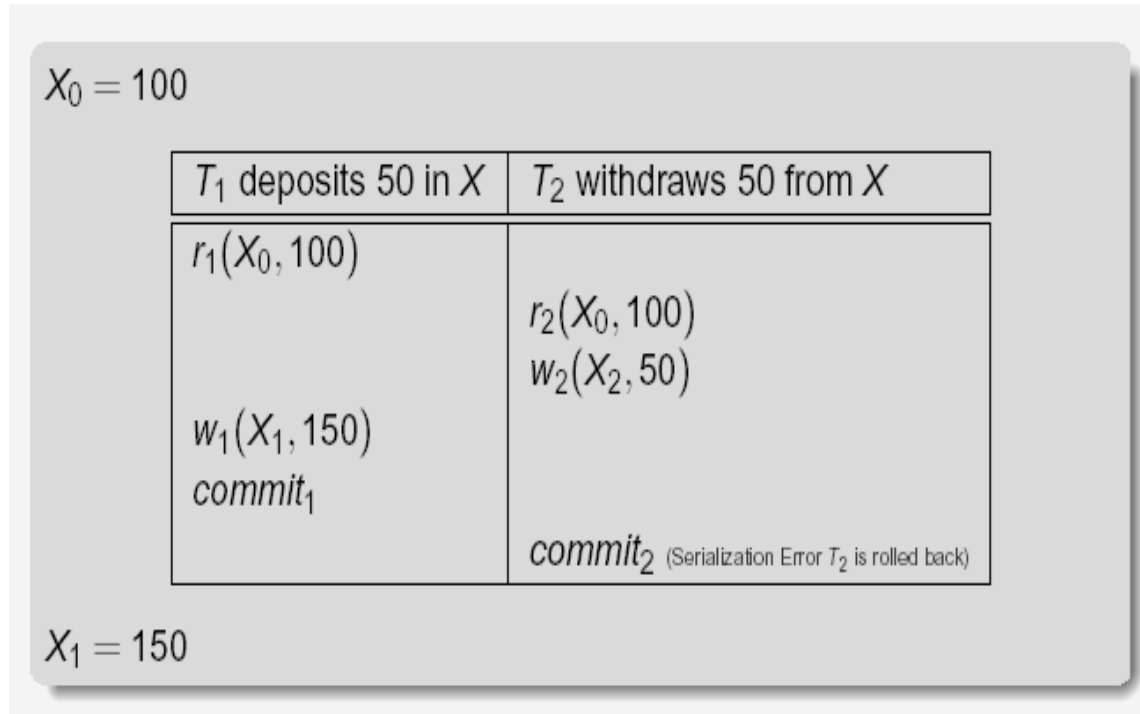
$X_0 = 100, Y_0 = 0$

T_1 deposits 50 in Y	T_2 withdraws 50 from X
$r_1(X_0, 100)$ $r_1(Y_0, 0)$ $w_1(Y_1, 50)$ $r_1(X_0, 100)$ (update by T_2 not seen) $r_1(Y_1, 50)$ (can see its own updates)	$r_2(Y_0, 0)$ $r_2(X_0, 100)$ $w_2(X_2, 50)$ $r_2(Y_0, 0)$ (update by T_1 not seen)

$X_2 = 50, Y_1 = 50$



Snapshot Write: First Committer Wins



- Variant: “**First-updater-wins**”
 - Check for concurrent updates when write occurs by locking item
 - ▶ But lock should be held till all concurrent transactions have finished
 - (Oracle uses **this plus some extra features**)
 - Differs **only in when abort occurs**, otherwise equivalent



Benefits of SI

- Reads are *never* blocked,
 - and also don't block other txns activities
- Performance similar to **Read Committed**
- Avoids **several anomalies**
 - **No** dirty read, i.e. no read of **uncommitted data**
 - **No lost** update
 - I.e., update made by a transaction is overwritten by another transaction that did not see the update)
 - No non-repeatable read
 - I.e., if read is executed again, it will see the same value
- Problems with SI
 - SI does not always give **serializable** executions
 - Serializable: among two concurrent txns, one sees the effects of the other
 - In SI: neither sees the effects of the other
 - Result: Integrity constraints can be violated



Snapshot Isolation

- Example of problem with SI
 - Initially $A = 3$ and $B = 17$
 - Serial execution: $A = ??$, $B = ??$
 - if both transactions start at the same time, with snapshot isolation: $A = ??$, $B = ??$
- Called **skew write**
- Skew also occurs with inserts
 - E.g:
 - Find max order number among all orders
 - Create a new order with order number = previous max + 1
 - Two transaction can both create order with same number
 - Is an example of phantom phenomenon

T_i	T_j
read(A)	read(A) read(B)
read(B)	
$A=B$	$B=A$ write(B)
write(A)	

