

# Architecture of Large-Scale Systems

By Dr. Taghinezhad

Mail:

a0taghinezhad@gmail.com

<https://ataghinezhad.github.io/>

# **CHAPTER 2 Two Mistakes High—Having Room to Recover from Mistakes**

# The Importance of Redundancy and Availability

- a MySQL database backup replica being used for experimentation, leading to a failure during a primary database outage.
  - Problem, Can backup replica be used for production, when it is experimented ?
    - No, because its setting is changed and it is not longer reliable
- **Key Lesson:** Backup systems must be treated with the same rigor as primary systems to ensure availability.
- **Takeaway:** Redundant systems are not just backups; they are critical components in maintaining high availability

## Two Mistakes High: Philosophy of Recovery

- **Concept:** "Two mistakes high" from radio-controlled planes.
- **Key Idea:** Always keep enough "altitude" (resources) to recover from two independent mistakes.
- **Application:** In highly available systems, plan for multiple failures and ensure recovery from any combination of mistakes.

<https://ataghinezhad.github.io/>

# Scenario #1: Node Failure

- **Initial Setup:** Service designed to handle 1,000 req/sec with how many number of nodes that each handles 300 req/sec each.
- Question: How many nodes do you need to handle your traffic demands? Some basic math should come up with a good answer:
  - $number\_of\_nodes\_needed = \lceil \frac{number\_of\_requests}{requests\_per\_node} \rceil$
  - can you handle the expected traffic, and because you have four nodes, you can handle the loss of a node?




Figure 2-1. Four nodes, 250 req/sec each

# Scenario #1: One Node Failure

- **Failure Situation:** One node fails; remaining nodes overloaded, leading to service degradation (Figure 2-2).
  - $\text{requests\_per\_node} = 1,000 \text{ req/sec} / 3 \text{ nodes} = 333 \text{ req/sec/node}$
  - That's 333 req/sec per node, which is well above your 300 req/sec node limit (see Figure 2-2).
- **Solution:** Add a 5th node to ensure capacity even after one failure (Figure 2-3).




Figure 2-3. Five nodes; one failure can still be handled





Figure 2-2. Four nodes; one failure causes overflow

## Scenario #2: Rolling Upgrades and Node Failures

- **Upgrade Plan:** Rolling deploy with 5 nodes ensures availability during upgrades.
- **Risk:** A node failure during an upgrade leaves only 3 nodes handling traffic, leading to an outage.
- **Lesson:** Ensure redundancy covers both routine maintenance and unexpected failures.




*Figure 2-4. Four data centers, 40 nodes, sufficient capacity to handle load*

## Scenario #2: Rolling Upgrades and Node Failures

- Suppose that you have a service whose average traffic is 1,000 req/sec.
- let's assume that a single node in your service can handle 300 req/sec.
  - Four node is enough to handle expected traffic
- You want to do a software upgrade while running your service nodes.
  - A rolling deploy (upgrade nodes one by one to keep operational reset when one is upgrading).
  - **How many nodes is needed?**
  - Five Nodes
    - This system can tolerate single node failure and support rolling deploy updates
    - Six Nodes, can handle multimode failure

# Scenario #3: Data Center Resiliency

- **Setup:** Service requires to handle 10,000 req/sec,
  - It would need 34 nodes without considering redundancy for failures.
  - let's uses 40 nodes across four data centers so that we have even more redundancy and fault tolerance.
  - Are we resilient?



*Figure 2-4. Four data centers, 40 nodes, sufficient capacity to handle load*

## Scenario #3: Data Center Resiliency (Cont.)

- Risk:** One data center outage leads to overloading the remaining nodes (Figure 2-5).




Figure 2-5. Four data centers, one failed, 30 nodes, insufficient capacity to handle load

- How many servers we need to handle lose of a datacenter?

$$\text{nodes\_per\_data\_center} = \lceil \text{min\_number\_of\_servers} / (\text{number\_of\_data\_centers} - 1) \rceil$$

- Solution:** To maintain capacity with a data center outage, 48 nodes are needed.

- $\text{nodes\_per\_data\_center} = \left\lceil \frac{34}{(4 - 1)} \right\rceil = 12 \text{ server/data_center}$
- How many nodes will it be?
  - $\text{total\_nodes} = \text{nodes\_per\_data\_center} \times 4 = 48 \text{ nodes}$

## Scenario #4: Hidden Shared Failure Types

- Sometimes seemingly independent problem scenarios can actually be dependent, meaning they might fail together
- **Example:** Your service needs four nodes, but you've wisely prepared with six nodes; enough to handle a single node failure and an upgrade in progress.
  - Six nodes sharing the same rack and power supply all fail simultaneously.
- **Key Point:** Ensure physical and infrastructure-level separation to prevent cascading failures.

# Scenario #5: Failure Loops

- A failure loop occurs when a problem prevents you from fixing it without causing a worse issue
  - **Example:** Imagine having a backup generator stored in your garage, but the only way to access the garage is through an electric-powered door that doesn't work during a power outage. Similarly, in the world of services, dependencies between failures and solutions can impact availability.
  - **Lesson:** Ensure that backup systems can be activated even during failures, avoiding failure loops.

# Managing Your Applications

- **Key Principles to manage your applications:**

- **“Fly Two Mistakes High”**
  - Look beyond surface failure modes.
  - Consider dependent failure layers.
  - Ensure recovery mechanisms work during failures.
- **Don’t Ignore Problems**
  - Persistent issues affect availability plans.
  - Backup systems matter—treat them seriously.
- **Production Is Production**
  - Everything in production matters.
  - Backup databases are mission-critical too.
- **Layered Failures Are Tricky**
  - Identifying dependencies isn’t obvious.
  - Invest time in understanding and resolving.

## Case Study: Space Shuttle Redundancy

- The **Space Shuttle software system** was one of the first large-scale applications to implement extreme redundancy and failure management.
  - Primary system: **5 computers** (4 identical running the same software, 1 independent).
  - Main process on all computes during critical parts:
    - **4 computers received the same data** and performed the same calculations.
      - If one computer differed, it was voted out and shut down as it was uncorrect. (**winners rule, loses terminate**)
        - The shuttle could operate with 3 computers and land safely with 2.

<https://taghinezhad.github.io/>

# Case Study: Space Shuttle Redundancy

- The **Space Shuttle's computer system** used redundancy to ensure safety and reliability:
  - **4 identical computers** ran the same software and performed identical calculations. If one failed, the others continued without interruption.
  - **1 independent backup computer** ran different software to double-check the main system's results.
- During **critical operations** like launch and landing, all four main computers processed the same data simultaneously.  
The shuttle could safely **operate with three** and even **land with two** working computers.
- This redundancy-based design ensured **30 years of missions** without any **life-threatening software failures** — a remarkable achievement in aerospace engineering.

<https://ataghinezhad.github.io/>

- End of Chapter 2

<https://ataghinezhad.github.io/>